

# Football Broadcast Player Detection and Re-Identification using YOLOv11

Final Project Report

Arpit Chauhan

## 1 Objective

The objective of this project is to design a robust, real-time system capable of detecting and consistently tracking football players in broadcast videos. A primary challenge addressed is ensuring that each player maintains a unique and persistent identity across video frames—even when temporarily occluded or out of view due to camera motion or scene transitions.

## 2 Approach and Methodology

The following steps were undertaken to build the Football Player Detection and Re-Identification System:

1. **Frame Extraction:** A 15-second football broadcast video was used as the input source. A total of 75 frames were extracted using OpenCV for annotation and training.
2. **Dataset Creation using Roboflow:** Each frame was uploaded to Roboflow, where manual annotations were performed. The labeled categories included:
  - Football players
  - Referee
  - Football
  - Other objects (e.g., goalposts, etc.)

Data augmentation techniques such as flipping, brightness/contrast adjustment, and blurring were applied. The final dataset contained:

- 238 images total
- 222 images for training
- 11 images for validation
- 5 images for testing

The dataset is publicly available on Roboflow at: <https://app.roboflow.com/arpittchauhan/devpmbm/4>

3. **Model Selection and Training:** The YOLOv11 model was selected for object detection. The training process was carried out using the Jupyter Notebook file `Football_Analysis_System (1).ipynb`, where the model was trained for 150 epochs with early stopping applied. The best model checkpoint was saved as `best.pt`.
4. **Post-Training Inference:** Two inference pipelines were implemented:
  - `yolo_inference.py` – for basic object detection using pretrained weights
  - `main.py` – for detection and player tracking using the custom-trained model
5. **Re-Identification and Output:** The output included annotated videos with consistent player IDs, and tracking metadata was saved as a `.pkl` file inside the `tracker_stubs/` folder.

### 3 Techniques Tried and Outcomes

- **Pretrained YOLOv11 Inference:** Initially, the `yolo_inference.py` script was used to perform basic object detection using pretrained YOLOv11 weights. The model was able to detect generic classes (players, ball) but without persistent IDs or tracking across frames.
- **Custom YOLOv11 Training and Re-ID:** A custom dataset was prepared and annotated using Roboflow. The YOLOv11 model was trained for 150 epochs. The training achieved a final mAP@50 of **0.888** and mAP@50–95 of **0.429**, with high precision and recall across classes like `football_players`, `referee`, and `football`.
- **Custom Tracking with Re-Identification:** The `main.py` script implemented a full pipeline using a custom tracker. The model was initialized with the fine-tuned weights, and the player tracking results were written to `tracker_stubs/player_detection.pkl`. Visual annotations were drawn across all video frames, and a final output video was saved in `output_videos/`.

### 4 Challenges Encountered

- **Re-ID Gaps After Frame Exit:** The model struggled to reassign the same identity to players who left the field of view and returned after a few seconds, leading to fragmented tracking.
- **Visual Similarity Confusion:** Players with nearly identical uniforms, especially teammates, caused the model to incorrectly swap or duplicate identities.
- **Blurred or Occluded Frames:** Fast movements or partial occlusions led to degraded visual features, making identity re-matching highly unreliable.

## 5 Model Performance and Results

The YOLOv11 model was trained for a maximum of 150 epochs with early stopping enabled (patience = 100). Training was halted at epoch 148, with best performance at epoch 48.

### Final Evaluation Metrics (Validation Set)

Category	Precision	Recall	mAP@50	mAP@50-95
Football	0.811	0.778	0.747	0.373
Football Players	0.984	0.922	0.953	0.469
Referee	0.918	0.929	0.963	0.445
Overall	0.904	0.876	0.888	0.429

Table 1: Per-Class Detection and Re-Identification Performance

### Speed Performance

- Preprocessing Time: 0.2 ms/image
- Inference Time: 17.4 ms/image
- Postprocessing Time: 1.1 ms/image

```

all      11      152      0.87      0.883      0.868      0.373

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
147/150  10.4G    0.8305    0.3913    1.048     156        640: 100% 14/14 [00:10<00:00, 1.39it/s]
Class   Images  Instances  Box(P   R       mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.31it/s]
all      11      152      0.872    0.883    0.864    0.371

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
148/150  10.4G    0.8166    0.3901    1.045     170        640: 100% 14/14 [00:10<00:00, 1.38it/s]
Class   Images  Instances  Box(P   R       mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.69it/s]
all      11      152      0.876    0.878    0.86     0.364

EarlyStopping: Training stopped early as no improvement observed in last 100 epochs. Best results observed at epoch 48, best
model saved as best.pt.
To update EarlyStopping(patience=100) pass a new patience value, i.e. `patience=300` or use `patience=0` to disable EarlyStop
ping.

148 epochs completed in 0.492 hours.
Optimizer stripped from runs/detect/train2/weights/last.pt, 51.2MB
Optimizer stripped from runs/detect/train2/weights/best.pt, 51.2MB

Validating runs/detect/train2/weights/best.pt...
Ultralytics 8.3.160 Python-3.11.13 torch-2.6.0+cu124 CUDA:0 (Tesla T4, 15095MiB)
YOLO11l summary (fused): 190 layers, 25,282,396 parameters, 0 gradients, 86.6 GFLOPs
Class   Images  Instances  Box(P   R       mAP50   mAP50-95): 100% 1/1 [00:00<00:00, 4.16it/s]
all      11      152      0.904    0.876    0.888    0.429
football 9        9        0.811    0.778    0.747    0.373
football players 11     131     0.984    0.922    0.953    0.469
referee  7        12       0.918    0.929    0.963    0.445

Speed: 0.2ms preprocess, 17.4ms inference, 0.0ms loss, 1.1ms postprocess per image
Results saved to runs/detect/train2
🔗 Learn more at https://docs.ultralytics.com/modes/train

```

Figure 1: YOLOv11 Training Summary: Loss curves, class-wise metrics, and evaluation output

## 6 Future Work and Possible Extensions

To enhance performance and robustness, especially in complex real-world settings, the following improvements can be considered:

- **Strengthen Re-Identification Across Gaps:** Train the system to better associate players who momentarily exit and re-enter the frame using more refined embeddings or temporal matching techniques.
- **Develop Player-Specific Re-ID Features:** Tailoring Re-ID by using features like jersey text, body build, or visual patches like logos could make the system more resilient to appearance overlap.
- **Apply Motion-Based Forecasting:** Motion prediction strategies such as Kalman filters or trajectory estimation can help maintain consistent identities during short disappearances.
- **Boost Efficiency for Real-Time Applications:** Convert and optimize the model using inference engines such as ONNX or TensorRT to reduce frame processing delays.
- **Scale to Multi-View Tracking:** Extend the architecture to handle multi-camera input and match players across views using cross-camera Re-ID techniques.

## 7 Conclusion

This project successfully demonstrated a working pipeline for football player detection and re-identification in broadcast videos using YOLOv11. The system showed significantly better performance when trained on a custom dataset and integrated with a tracking mechanism. Future improvements can include real-time deployment, jersey number OCR, and larger-scale datasets for robust generalization.